<u>Analysis of Issues</u>

AI Scraping and the Future of Digital Preservation Archives

Hrishika Roychoudhury (PID: 730557151)

INLS 690: AI Law and the Struggle for Control of Knowledge

Professors Will Cross and Dave Hansen

April 2026

Table of Contents

What the Law Requires

The law gives preservation institutions both a floor and a ceiling. Neither permits the easy solution of treating all AI access as either harmless or forbidden.

Copyright and fair use. The reproduction right under 17 U.S.C. §106 attaches the moment a crawl copies a protected page. Every bulk extraction of archived content is prima facie infringement unless a defense applies. The principal defense available to computational users is fair use under §107, and the relevant case law now draws a workable line, even if courts have not yet applied it directly to digital archives.

Authors Guild v. HathiTrust, 755 F.3d 87 (2d Cir. 2014) established that digitizing millions of books to build a searchable index was transformative fair use because the purpose (search and accessibility) differed fundamentally from the original purpose (reading). The court emphasized that the copies produced by the university libraries were “not commercially exploitative” and that “the libraries did not seek any commercial advantage from their use.” The court also confirmed that preservation copies for disaster recovery are lawful under §108. *Authors Guild v. Google* extended this reasoning to Google Books' snippet display. These holdings protect the Archive's own preservation activity and most scholarly text mining. However, the moment a commercial AI company trains a general-purpose foundation model on the same archives and sells that model as a service, the HathiTrust analysis becomes considerably more complicated.

There is a line drawn in *Thomson Reuters v. Ross Intelligence*, No. 1:20-cv-613 (D. Del. Feb. 11, 2025) with commercial use. The court held that Ross's use of Westlaw headnotes to train a competing legal research tool was not fair use because Ross used the material in precisely the way it was designed to be used, and the resulting product competed directly with the original. The court applied all four §107 factors and found three favored Thomson Reuters, illustrating the limiting principle: the fair use defense fails when the AI tool's outputs function as a 'market substitute' for the original copyrighted work. Under that standard, a model trained on archived journalism that then summarizes news for paying customers may not be protected, while the same model used to study journalistic language patterns likely is. The U.S. Copyright Office's May 2025 report on generative AI training reaches a compatible conclusion through different reasoning, placing "noncommercial research or analysis that do not enable portions of the works to be reproduced in the outputs" at the clearly fair end of the spectrum, and "copying of expressive works from pirate sources in order to generate unrestricted content that competes in the marketplace, when licensing is reasonably available" at the clearly infringing end. Most real-world AI uses fall between those poles.

The court in *Bartz v. Anthropic*, No. 3:24-cv-05417-WHA (N.D. Cal., settlement prelim. approved Oct. 17, 2025), found that Anthropic's reproduction of named plaintiffs' books in the process of training its models was fair use, but found that downloading millions of books from pirate shadow libraries to build its central collection was not fair use--at least not on the summary-judgment record. That distinction, between how AI companies use what they have and how they obtained it in the first place, runs directly through the archive context. Even if training on a public archive would otherwise be fair use, stripping CMI, exceeding rate limits, or obtaining access through deliberate circumvention of technical controls changes the analysis.

Kadrey v. Meta Platforms, 2025 WL 1786418 (N.D. Cal. June 27, 2025), and related cases have begun to test whether CMI removal during AI training can be independently actionable even where the underlying training might otherwise qualify as fair use. Because Meta's training copying was deemed fair use, the plaintiffs' DMCA §1202(b)(1) claim for removal of copyright management information also failed. The court's reasoning clearly stated that Congress could not have intended to exempt fair uses from copyright liability while simultaneously exposing the same users to DMCA liability for incidental CMI removal. The legal "must" here is narrow but important. If training is fair use and CMI was removed as part of standard pre-training data cleaning, the DMCA claim is unlikely to succeed. If training is not fair use, CMI removal becomes an independent ground for liability.

What the law does not protect. Section 230 of the Communications Decency Act does not help archives facing copyright claims. The statute explicitly states "Nothing in this section shall be construed to limit or expand any law pertaining to intellectual property." 47 U.S.C. §230(e)(2). Nor does the Computer Fraud and Abuse Act (CFAA) provide a useful tool for archives trying to block scrapers. *hiQ Labs v. LinkedIn*, 31 F.4th 1180 (9th Cir. 2022), and *X Corp. v. Bright Data*, No. C 23-03301 JSW (N.D. Cal. Apr. 1, 2024), both hold that accessing publicly available data is not "unauthorized access" under 18 U.S.C. §1030. An archive cannot threaten scrapers with criminal liability for collecting material the archive itself has made publicly available.

However, although these cases conclude that public data alone does not give rise to a CFAA claim, the analysis changes when access happens at a scale and through methods that exceed any reasonable reading of authorization. The 2023 conduct, with rotating cloud hosts and continued requests after explicit blocking, sits closer to a denial-of-service event than to ordinary scraping. The Archive must publish clear bulk access conditions and treat their violation as the basis for revoked authorization. The Archive may not, however, frame normal access as unauthorized; the law closes that path with 18 U.S.C. §1030.

What contract can do. *ProCD, Inc. v. Zeidenberg*, 86 F.3d 1447 (7th Cir. 1996), establishes that contract conditions on the use of information can be enforced even when copyright cannot reach the same result, because contracts bind their parties rather than the world at large. A Terms of Service agreement, a Data Use Agreement, or a tiered API license can impose reciprocity requirements, prohibit certain categories of use, and authorize termination for violations. The Archive may not be able to sue a company for scraping under copyright, but it can enforce an agreement that the company signed to obtain structured API access. This is the legal mechanism that makes a credentialing framework operable.

The Archive's existing browse-wrap terms are weak under the analysis Lemley and Henderson set out. The Archive must, if it wants its bulk access conditions to bind, move to a model that requires registration, agreement, and an issued credential before high-volume access is granted. It may not rely on terms buried in a footer to constrain anonymous scraping.

On Section 1201, the Archive may deploy authentication, rate limits, anomaly detection, and similar technical access controls, and circumvention of those controls becomes a statutory violation independent of any underlying copyright infringement. The 1201 rulemaking has carved out exemptions for nonprofit text and data mining; the policy must respect those exemptions and not block bona fide research. On Section 1202, the Archive should require that bulk users preserve copyright management information rather than strip it, since metadata stripping is now a contested issue in *Kadrey v. Meta* and related cases.

Where the Law leaves Discretion

Beyond the problems that the law decides on, are the ones it leaves to the institution. Several questions the policy must answer are not answered by statute or case law. These require institutional choices, and each choice should be explained to stakeholders rather than dressed up as legal compulsion.

The first is whether to treat AI training as categorically different from scholarly text and data mining. The Library Copyright Alliance's 2023 principles state plainly that the ingestion of copyrighted works to create AI training databases is generally a fair use, and that the contribution of any one work to a model is *de minimis*/insignificant. But the technical reality is that a graduate student running a corpus analysis and a commercial AI lab building a foundation model use indistinguishable methods. The classification cannot be made at the network level.

The retrieval-versus-training distinction. The UVA Archival AI Protocol draws a "sharp line" between retrieval-based AI, which searches a collection at query time while the institution retains control, and training-based AI, which absorbs materials into model weights in a process that cannot be undone. The law does not mandate this distinction. HathiTrust approved transformative non-expressive copying regardless of the technology involved. But the distinction is sound policy because it tracks institutional control and the reversibility of decisions. A policy that defaults to allowing retrieval-based uses and requiring formal review for training-based uses reflects the Archive's actual relationship to its materials, not a legal technicality.

Technical controls occupy the most immediate and tractable policy space. NC State University Libraries and similar institutions have implemented tiered network access systems in which users connected to institutional networks receive priority access, external academic users receive standard access, and automated bots can be identified and throttled or blocked at the network level. The Archive's open, internet-accessible infrastructure does not naturally support such tiering, but adapted versions are possible.

Preference signals. The Library Copyright Alliance Principles state that "copyright owners can employ technical means such as the Robots Exclusion Protocol to prevent their works from being used to train AIs." That is a statement about what rights holders may do. The Budapest Open Access Initiative's 2002 framing explains the only appropriate constraint on reproduction is "to give authors control over the integrity of their work and the right to be properly acknowledged and cited."

Some other strategies I've seen implemented include: rate limiting by IP address and user-agent string; CAPTCHA-based human verification for access above defined request thresholds; robots.txt enforcement for named crawlers; and the emerging ai.txt standard developed by Spawning.ai, which allows websites to signal 'archiving permitted, AI training prohibited' at the bot level. The W3C is developing analogous TDM-specific machine-readable standards; the European Union's copyright directive already recognizes machine-readable TDM opt-outs as legally operative, and U.S. policy development is tracking this.

However, outside the European Union, robots.txt and the emerging ai.txt standard carry no legal enforcement weight in U.S. courts. They work only for bots that voluntarily respect them--and we have seen incidents involving bots that manifestly did not. A policy that relies solely on voluntary compliance with machine-readable signals is no policy at all for actors operating in bad faith. The Archive must therefore combine machine-readable signaling with active monitoring, IP blocking capability, and legal mechanisms for seeking injunctive relief against bad-faith extractors. None of these mechanisms is perfect in isolation; the point is that each adds friction and cost for bad actors while preserving access for compliant users.

The second question is whether to charge for commercial bulk access. The law does not require fees, but the Archive's nonprofit status, current financial pressure, and the fact that it has just settled a major dispute with music publishers over the Great 78s project all argue for a sustainable funding model. My proposed answer is yes, scaled to the commercial value of the use, with proceeds dedicated to preservation infrastructure. This is what Harvard Library's public domain corpus project did in a different form, and what the Library Copyright Alliance's principles, while skeptical of compulsory licensing, do not foreclose for voluntary commercial agreements. Brewster Kahle has proposed publicly that AI companies could return value to the Archive by implementing 'anti-hallucination' features that link chatbot assertions to verified archival citations, effectively using the Archive's collection of pre-generative AI content to anchor the reliability of AI outputs. This proposal reframes the relationship from extraction to partnership.

Reciprocity. The law does not require AI companies to contribute anything back to the institutions whose collections they use. The policy's reciprocity requirements, contributing derived datasets, providing transparency reports, or implementing citation-based anti-hallucination features as Kahle has proposed, rest entirely on contractual agreement. The justification for requiring them is practical however, as an institution that offers structured access to curated, high-quality historical data is offering something worth more than free bulk extraction. Value should flow in both directions. An institution whose infrastructure is strained by tens of thousands of requests per second, whose preservation mission is jeopardized by the publisher withdrawals those requests provoke, and whose operating budget depends on public and philanthropic support, has a legitimate institutional interest in requiring that commercial beneficiaries of its work contribute to its sustainability.

Other forms of reciprocity are available. AI companies that train on archival collections could contribute derived datasets back to the Archive, improving its collections and providing new computational tools for future researchers. They could provide usage transparency reports that document what was accessed, when, and for what purpose. They could agree to fund the Archive's preservation infrastructure at a level proportionate to the value extracted. The Harvard Library Public Domain Corpus model, in which Harvard made nearly one million digitized public domain books available with full documentation of provenance and terms, demonstrates that “documented, legally sound data” can “shift the power dynamic” back to public institutions rather than allowing private commercial entities to filter all access through proprietary systems. The Data Governance in Open Source AI white paper published by the Open Source Initiative in January 2025 frames this as a structural problem with the current AI data ecosystem, noting that "open sharing of various resources needs to go hand in hand with greater attention placed on data quality, data governance, responsible sharing and respect and protection for various rights in data."

The third and final question is whether to take a firm public position on the publisher withdrawal cascade. According to analysis published in early 2026 by the AI-detection startup Originality AI, 241 news sites in nine countries are now blocking at least one of the Internet Archive's four crawling bots, with about 87 percent of those sites owned by USA Today Co. The New York Times has gone further, employing what reporting describes as hard-blocking measures beyond standard robots.txt. Reddit announced in August 2025 that it would block the Archive after concluding that AI companies were using the Wayback Machine as a back door to content Reddit was licensing for revenue. EFF has compared publisher blocking of preservation crawlers to a newspaper publisher refusing to let libraries keep copies of its paper.

Privacy Law. Under the California Consumer Privacy Act and analogous state laws, the Archive may face deletion obligations for personal data held within its collections if residents of covered states make valid deletion requests. This is not a wholesale obligation to alter historical archives, but it does require that the Archive have a workable process for handling such requests, particularly where archived pages contain sensitive personal information whose subject never consented to archival.

The intersection with AI training is more directly consequential, as a model trained on archived personal data may retain and reproduce that data in outputs, creating downstream privacy harms the Archive cannot control once the training has occurred. The UVA Archival AI Protocol's 'right to stop' provision, which allows institutions to order the cessation of data ingestion and demand decommissioning of models built on defined collections if permission is withdrawn, is a direct response to this problem. The Archive should adopt analogous contract language in any formal access agreements.

My further proposed answer is that the Archive should publish a clear position acknowledging the publishers' AI concerns as legitimate, while pointing out that blocking preservation does not solve the AI problem. The Guardian's surgical approach, in which it limits Archive access to its journalism temporarily while continuing to permit archiving of its regional homepages and topic pages, suggests a middle path the Archive might want to encourage and document. As Mark Graham has put it, the Archive is "collateral damage in a copyright war it didn't start." The best institutional response seems to be to stop being collateral damage by giving publishers a credible structure for distinguishing preservation from extraction. The Archive should position itself, in its communications with publishers, as a potential ally rather than a passive conduit. If the Archive can credibly demonstrate that it will block AI extraction from archived content at the publisher's request while still preserving historical versions for non-commercial research and journalism, the case for hard blocking weakens substantially.

Balancing Stakeholder Needs

Researchers need broad access, including computational access. The policy protects them by creating a research tier rather than forcing them into the same bucket as commercial AI companies.

Journalists need fast, public access to archived pages for accountability reporting. The policy keeps ordinary Wayback access open and resists publisher hard blocks.

Publishers need protection against AI companies using archives to bypass licensing. The policy gives them granular signals, commercial-use restrictions, and a path to data agreements.

AI developers need lawful, structured, high-quality data. The policy gives them a front door instead of forcing them into scraping or litigation.

Archivists and IT staff need a system they can run. The policy uses rate limits, APIs, logs, and review tiers rather than case-by-case crisis response.

Communities represented in collections need safety, context, and control. The policy creates sensitive-data review and does not treat openness as the only value.

Drafting Policy Language

There are three principles that should be practiced as I outline a framework for the policy to follow.

The first is to state conditions as if/then rules rather than values. "We support open access" is a value statement that no one can apply. "Any request to access the API at more than 1,000 requests per hour requires submission of a Data Management Plan and execution of a Data Use Agreement before access is granted" is a rule that can be applied by anyone. The policy will use a tiered access framework that mirrors the structure of the University of Virginia Archival AI Protocol, which Leo S. Lo published through the UVA Library in January 2026 and which a number of archival organizations have begun adopting. The UVA approach distinguishes retrieval-based AI uses, which leave source materials under institutional control and are presumptively allowed, from training-based uses that absorb materials into model weights and are presumptively restricted unless provenance, attribution, and a real right to stop further use can be guaranteed.

The second is to sort uses into categories before writing permissions for each. The five scenarios from the Scenario Analysis, a journalist using Wayback snapshots, a graduate student running corpus analysis, a research team running large-scale NLP, an AI safety lab building a fact-checking tool, and a commercial company seeking bulk extraction for a general-purpose model, represent a real range of uses the policy must address. Language that does not produce different outcomes for the commercial extraction case and the graduate student case has failed.

The Archive's adaptation has three tiers. Tier one is unauthenticated public browsing of the Wayback Machine and the rest of archive.org. It is unconditional and remains the heart of the institution's mission. Tier two is a registered Researcher Passport for nonprofit, scholarly, and journalistic bulk access through the Scraping API, with a brief data management plan, rate limits calibrated to legitimate research use, attribution requirements, and prohibition on commercial AI deployment. The Passport draws on the ICPSR restricted data access framework provided in their white paper series to strategize for a Community-normed System of Digital Identities of Access. This tier may need to be reconsidered if the Archive finds implementing this barrier to be too far against their fundamental open-access principles. Tier three is a Commercial Bulk Access License for everyone else, with a fee, transparency reporting, machine-readable preference signal compliance, copyright management information preservation, a contractual right to stop further use, and a reciprocity requirement that takes one of several forms including the anti-hallucination citation linkage Brewster Kahle publicly proposed.

The third practice is to define every term that will bear legal weight. "Commercial purpose," "bulk extraction," "high-volume access," "retrieval-based," and "training-based" all need definitions in the policy document itself. Undefined terms invite the disputes that policies are meant to prevent. The Copyright Office's May 2025 report is a good model here, defining "pre-training," "fine-tuning," "RAG," and "memorization" with enough precision that the fair-use analysis can be applied differently to each stage. The NASCIO-CoSA State Archiving Playbook (2018) also offers a useful precedent from the records management world; its 'plays' are concrete, action-oriented, and specific enough to be implemented by non-lawyers. The Archive's AI access policy should adopt a similar structure, distinguishing clearly between categories of permitted use, uses requiring credentialing, and prohibited uses.

Permitted use, without credentialing: any individual access, research citation, historical inquiry, journalism, or litigation support that does not involve automated bulk extraction. Credentialed use: bulk extraction for documented academic or public-interest research, subject to data management plan, data use agreement, CMI preservation requirement, and right-to-stop clause. Prohibited use: bulk extraction for commercial AI training or commercial product development without a formal Data Trust Agreement that includes reciprocity obligations, provenance standards, and the right to stop and decommission.

Schafer and Winters's framework for web archive governance, with its emphasis on participation, accountability, transparency, effectiveness, inclusivity, and legality, provides a useful checklist for evaluating draft policy language. A policy that meets all seven criteria will be more durable and more defensible than one optimized for a single value. The goal is a framework that allows the Archive to say, credibly and publicly, that it has done everything within its power to protect publishers' legitimate interests, support legitimate research, and resist extractive commercial exploitation, while remaining the institution that democracy and public memory require it to be.

The analysis that follows in the Final Annotated Draft Policy will apply these three practices to produce specific, testable language for each tier of the access framework.