

Final Annotated Draft Policy

A Framework for the Internet Archive's Access and Use Policy for Large-Scale Automated Requests

INLS 690: AI Law and the Struggle for Control of Knowledge

Professors Will Cross and Dave Hansen

Table of Contents

[Table of Contents](#)

[I. The Policy and Its Context](#)

[II. The Problem and Why a Policy is Necessary](#)

[III. Stakeholder Perspectives](#)

[IV. Obstacles and Challenges](#)

[A. Legal Uncertainty](#)

[B. Technical Limitations](#)

[C. Institutional Capacity](#)

[V. Decision Points and Proposed Resolutions](#)

[VI. Annotated Path Forward](#)

[VII. Sources Cited](#)

TO: Internet Archive Board of Directors, Office of General Counsel, and Senior Leadership

FROM: Hrishika Roychoudhury, Graduate Information Science Student

DATE: May 2026

RE: On AI Scraping and the Future of Digital Preservation Archives

I. The Policy and Its Context

The following proposed policy text is designed to govern bulk automated access to the Internet Archive's digital collections, including those accessible through the Wayback Machine. It would be incorporated into the Archive's existing Terms of Service as a new section, supplemented by an internal implementation guide for Archive staff, a model Data Trust Agreement for the legal team, and a public-facing FAQ for researchers and developers. The policy text is reproduced first, followed by the analysis that justifies each provision and annotations explaining the reasoning behind each section.

PROPOSED POLICY TEXT: AUTOMATED AND BULK ACCESS TO INTERNET ARCHIVE COLLECTIONS

***Section 1. Purpose and Scope.** This policy governs any access to the Internet Archive's collections that involves automated requests at rates exceeding 100 requests per second,*

any programmatic bulk extraction of content for training machine learning models or artificial intelligence systems, and any use of archived content for commercial AI product development. It does not apply to individual access, standard research browsing, journalistic citation, litigation support citation, or academic citation at normal human interaction speeds.

Section 2. Open Access Tier. *Any user who accesses Archive collections at rates below the threshold established in Section 1, for any purpose, requires no registration. The Archive affirms its commitment to preserving and providing universal access to the historical record for individual users, researchers, journalists, educators, and the public.*

Section 3. Credentialed Research Tier *(Academic and Public-Interest Bulk Access).* *Researchers, including affiliated academic researchers, independent scholars, public-interest journalism organizations, community archivists, and students, who require bulk automated access for documented research purposes may apply for Credentialed Research Access. Applicants must (a) submit a Data Management Plan specifying the collection or collections to be accessed, the volume of data required, the research purpose, and the intended use of outputs; (b) agree to a Data Use Agreement prohibiting use of the data outside the described project, prohibiting attempts to re-identify individuals, and requiring preservation of Copyright Management Information in any derived outputs; (c) acknowledge the Archive's right to audit compliance and to require the cessation of data ingestion if the terms are violated. Credentialed Research Access will be granted or denied within thirty business days of a complete application. The Tier explicitly includes non-university pathways for independent researchers, journalists, local historians, and community archivists, so that the credentialing requirement does not function as exclusion.*

Section 4. Data Trust Tier *(Commercial and Foundation Model Training).* *Any person or entity seeking to use Archive collections for the purpose of training a commercial AI model, a general-purpose large language model, or any AI system whose outputs are monetized or provided as a commercial service must enter into a Data Trust Agreement with the Internet Archive before accessing collections in bulk. Data Trust Agreements require:*

- (a) a detailed Provenance and Attribution Plan demonstrating that the entity can trace AI outputs back to the specific archived items on which they were based, and that minimum provenance standards as specified in the Archive's published guidelines will be met;*
- (b) meaningful CMI preservation and attribution in all derived outputs;*
- (c) a reciprocity commitment in one or more of the following forms: a financial contribution to the Archive's preservation infrastructure proportionate to the volume and*

commercial value of the data accessed; the contribution of derived datasets back to the Archive under an open license; preservation services such as mirroring; or the implementation of citation and anti-hallucination features that link AI system outputs to verified archival sources;

(d) a binding right-to-stop clause under which the Archive may order the cessation of data ingestion and, for narrow-scope models trained mainly on defined Archive collections, may require decommissioning, if the agreement is violated. For large mixed-corpus foundation models, the right-to-stop is calibrated to cessation of new training, removal of Archive content from retrieval indexes, and documentation of past use, recognizing that decommissioning is not realistic;

(e) annual compliance reporting to the Archive, including disclosure of collections accessed, volumes extracted, intended and actual uses, and any commercial products incorporating Archive-derived material;

(f) a contractual prohibition on producing outputs that directly substitute for original publishers' core commercial products without separate licensing arrangements with those publishers, allocating the market-substitution risk addressed in Thomson Reuters v. Ross to licensees rather than to the Archive.

Section 5. Machine-Readable Signaling. *The Archive will implement and maintain ai.txt files consistent with current and emerging machine-readable standards for signaling content permissions. Where content within Archive collections has been contributed by third-party publishers who have designated archiving-permitted but AI-training-prohibited status through machine-readable signals, the Archive will honor those designations in its bulk access agreements and will make good-faith efforts to implement technical enforcement.*

Section 6. Rate Limits and Infrastructure Protection. *Any automated access, regardless of tier, is subject to rate limits as published by the Archive in its technical documentation. Exceeding those limits without prior written authorization constitutes a violation of this policy. The Archive reserves the right to block IP addresses, virtual host ranges, or user-agent strings that violate rate limits or that the Archive has reasonable cause to believe are operating in bad faith, without prior notice. The Archive will implement application-layer authentication, rate limiting with graduated throttling, anomaly detection for distributed coordinated access, and updated robots.txt and ai.txt directives consistent with the tiered framework. These measures are technological protection measures within the meaning of 17 U.S.C. § 1201.*

Section 7. Publisher-Requested Restrictions. *Where the Archive has received documented requests from original content publishers to restrict bulk automated access to*

archived versions of that publisher's content, the Archive will implement those restrictions in Data Trust and Credentialed Research agreements and will make good-faith technical efforts to enforce them. This provision does not limit individual or journalistic access to archived publisher content. The Archive will publish an annual summary of active publisher partnership arrangements.

Section 8. Enforcement and Remedies. *Violation of this policy subjects the user to immediate termination of access, forfeiture of any access agreement, and potential legal action. The Archive will pursue injunctive relief against users who cause infrastructure overload or who engage in bad-faith extraction. Tier 2 violations may result in Credentialed Research Access revocation and breach of contract claims. Tier 3 violations may result in Data Trust Agreement termination, breach of contract claims, and where applicable claims under 17 U.S.C. §§ 1201 and 1202 and 18 U.S.C. § 1030. The Archive will publish an annual access governance report describing Credentialed Research holders by category, active Data Trust counterparties, reciprocity contributions received, ai.txt compliance status, and enforcement actions.*

II. The Problem and Why a Policy is Necessary

In April 2026, USA Today published an investigation that reconstructed how the U.S. Immigration and Customs Enforcement (ICE) had quietly altered its public detention statistics over time. The reporters could only do this because the Internet Archive's Wayback Machine had preserved the earlier versions of ICE's pages before the agency changed them. There is, as Wayback Machine director Mark Graham noted in Wired, "a little irony" in this. USA Today Co., the publishing conglomerate formerly known as Gannett and operator of more than 200 outlets nationwide, actively bars the Wayback Machine from archiving its own work (Knibbs, 2026). The story illustrates exactly what the Wayback Machine exists to do, and exactly how fragile that capacity has become.

The Internet Archive, the Library of Congress's web archiving programs, and similar digital institutions preserve the web because the web does not preserve itself. Pages disappear, URLs break, publishers revise stories, agencies remove pages, companies shut down sites, and platforms change access rules. Pew research found that a quarter of webpages that existed at some point between 2013 and 2023 were inaccessible by October 2023, and 38 percent of 2013 webpages were unavailable a decade later (Pew Research Center, 2024). Operating on a simple democratic premise that the open web should remain open for the future, these institutions work to preserve the record of the open web so that journalists, researchers, educators, courts, and future generations can know what was said, changed, deleted, or hidden. Web archives make it possible to investigate government revisions, track changes in news

reporting, recover dead links, preserve community history, and resist the quiet disappearance of digital culture.

Nowadays, AI companies need enormous volumes of text, images, metadata, and historical material. Digital archives are attractive because they contain structured, time-stamped, high-value materials at scale: public domain books, OCR files, news pages, government documents, web snapshots, and records that may no longer exist on the live web. The same architecture that makes these archives valuable for preservation also makes them vulnerable to extraction.

For instance, the Internet Archive is the largest digital preservation nonprofit in the world, holding over a trillion captured webpages. Journalists use it to track editorial changes; litigants submit archived pages as evidence; historians study how political language has shifted. On Memorial Day weekend in 2023, as the Archive's founder and director Brewster Kahle documented in a public blog post, tens of thousands of requests per second for OCR text files arrived from sixty-four virtual hosts on Amazon Web Services, crashing the Archive for roughly an hour. Upon blocking those IP addresses, in just a few hours another sixty-four resumed the same behavior. Kahle wrote at the end, 'Those wanting to use our materials in bulk should start slowly, and ramp up. Also, if you are starting a large project please contact us at info@archive.org, we are here to help.' The incident shows the problem goes beyond copyright and into infrastructure, access governance, and institutional survival; revealing that the Archive's infrastructure, designed for public access, is unequipped to negotiate terms with an extractive industry operating at machine scale. An open archive can be treated as a free input pipeline by commercial actors who do not bear the costs of preservation.

That incident was just the visible event. The structural problem behind it has three parts. First, AI training is effectively irreversible. Once a model has incorporated text into its weights, no one can extract that influence on demand. The UVA Archival AI Protocol takes this fact as its foundation, applying what it calls a principle of reversibility: archivists do not do what cannot be undone. Second, the Archive's own infrastructure invites the very behavior that hurts it. The public Scraping API at archive.org/services/search/v1/scrape provides cursor-based bulk access with no commercial restrictions, and commercial scraping vendors openly market the Archive as a target. Automatio advertises "Intelligent Retry Logic" as a feature for getting around the Archive's rate limits. The open-source Wayback Machine Scraper on GitHub provides configurable concurrency settings explicitly for scraping historical snapshots at scale. Third, the publishers most affected by AI training have begun blocking the Archive's preservation crawlers as their only available defense.

The downstream consequences are now playing out in real time. According to an analysis published in early 2026 by the AI-detection startup Originality AI, 241 news sites across nine countries now explicitly block at least one of the Internet Archive's crawling bots, with 23 major news sites blocking ia_archiverbot entirely. The Guardian has adopted a more surgical approach, blocking archival access to its journalism specifically while leaving topic and homepage content available. The New York Times has gone further, implementing 'hard blocks' beyond what standard robots.txt conventions require. The Times declared that “Times content on the Internet Archive is being used by AI companies in violation of copyright law to directly compete with us.” Reddit announced its own block in August 2025, citing AI companies' use of the Wayback Machine as a back door to content Reddit was licensing to Google for a reported sixty million dollars per year (Joseph, 2025; Knibbs, 2026; Sircar, 2026). As Anisha Sircar reported in Forbes, USA Today Co. alone accounts for 87 percent of the blocking publishers, given its ownership of more than 200 media outlets under the former Gannett umbrella. The Guardian's director of business affairs and licensing told Wired the outlet had been in conversation with the Archive over 'concerns over potential misuse by AI companies of content sets crawled for preservation purposes.' A 2023 analysis of Google's C4 training dataset by the Washington Post confirmed that the Internet Archive was among the sources used to build both Google's T5 model and Meta's Llama models.

These publisher responses are rational from their individual perspectives but collectively catastrophic for the public record. If major news organizations stop permitting archival crawling, there will be no historical record of early twenty-first century digital journalism. Wayback Machine director Mark Graham has described the Archive as “collateral damage in a copyright war it didn't start.” The Internet Archive's own *Vanishing Culture: A Report on Our Fragile Cultural Record* names the cumulative effect of these pressures a digital memory hole (Messarra, Freeland, & Ziskina, Internet Archive Press, 2026). The harm is recursive. As publishers block the Archive, the public web becomes less archivable. As the Archive's collections become less complete, the marginal value of scraping the Archive directly increases for AI companies now denied access to the publisher sites. As that scraping pressure increases, publishers are given more reason to block. The dynamic accelerates itself. The Blackmeta distributed denial-of-service attack of October 2024, which took the Archive offline for four days, was a separate but parallel reminder that the institution itself sits in a precarious infrastructure position. Without the Wayback Machine functioning, the ability to track editorial changes, verify deleted claims, and conduct accountability journalism will belong exclusively to organizations with the resources to build their own private archives. The policy proposed here is designed to prevent that outcome by giving publishers a reason to remain in the archival ecosystem rather than abandon it.

This is the scenario a policy must address: a nonprofit preservation institution that is both a victim of AI extraction and, in the eyes of publishers, an accessory to it. The policy challenge is not simply to choose between open access and content protection. It is to build a framework that can distinguish between the unannounced midnight AWS blitzkrieg and the graduate student's dissertation corpus, between the commercial foundation model and the fact-checker's text-mining script, and to do so with legal precision, institutional credibility, and the technical capacity to enforce it.

III. Stakeholder Perspectives

On May 6, 2026, the Internet Archive published an episode of the Future Knowledge podcast, "Preserving the Web in the Age of AI," in which Wayback Machine director Mark Graham, technology writer Mike Masnick of TechDirt, and media lawyer Kendra Albert of Albert Sellers LLP discussed the effect of AI on archives. Graham's framing in that conversation directly states the issue: the Wayback Machine has become "collateral damage caught up in the conflict between AI companies and publishers" (Freeland, 2026). There are a few points from that conversation that deserve to sit on the record of this policy. First, Graham noted that very few news organizations have actually taken the measures the New York Times and USA Today Co. have taken, and that the vast majority of publishers in the world remain happy for their material to be archived; without them, the record would not exist. Second, Andrew Deck of Nieman Lab, in his Marketplace reporting, could not find any examples where any publisher had identified evidence that material from the Wayback Machine had in fact been used by AI companies; the panic is operating ahead of the evidence (Deck, 2026). Third, Graham also observed that the Internet Archive has already implemented a series of measures, in collaboration with several news organizations, to limit bulk access to material from particular sites, including disabling certain enumeration features, partnering with Cloudflare on rate limiting, and optimizing the Wayback Machine for human use rather than automated extraction. The institution has already begun the work this policy proposes to formalize.

Masnick characterized the publisher blocking as "an overreaction to a problem that probably isn't really a problem" and suggested it would be seen in retrospect as a serious mistake by the media organizations responsible. He framed the broader context as a general cultural backlash against AI in which less sophisticated sites are blocking everything because that is the only lever they understand, even where the cost falls on archives rather than on AI companies. Albert added the legal layer, explaining that web archiving sits across two distinct legal frames. The fair use frame, where archiving fits squarely within the line running from cached web copies through HathiTrust, is reasonably favorable. The web scraping frame, governed by the Computer Fraud and Abuse Act, by Section 1201, by trespass to chattels, and by terms of service, is much less favorable, because those bodies of law typically do not ask why the

access is occurring. As Albert put it, "in the fair use context, you get to talk about things like, hey, it's really good for general knowledge that folks can access archived versions of websites; in many conversations around web scraping, you're often much less focused on the question of why are you doing this?"

Albert's caution about overbroad responses is directly relevant to this policy: "websites, whether they're news publishers or not, should be careful about the degree to which we throw the baby out with the bathwater and say, well, actually some of these entities are behaving badly or doing things that we don't like with our content, whether they're entitled to legally or not, and therefore we're just going to take a broad stance across the board." The same principle cuts the other way. The Archive's policy must avoid the same overbreadth in reverse. A policy that treats every bulk request as a potential threat will exclude the journalist, the graduate student, and the public-interest researcher along with the commercial scraper. The framework here is designed precisely to draw the distinctions Albert says the law fails to draw on its own, and to draw them through institutional accountability rather than through technical exclusion.

Journalists who use the Wayback Machine for investigative work are the Archive's most vocal constituency. More than one hundred reporters, ranging from Rachel Maddow to independent journalists like Kat Tenbarge and Taylor Lorenz, signed a letter organized by the Electronic Frontier Foundation, Fight for the Future, and Public Knowledge in April 2026 in support of the Archive (EFF, FFF, & Public Knowledge, 2026). They named what we would lose if preservation continued to erode: the ability to track editorial changes, to access old fan sites and cultural ephemera, to use the Archive in union organizing. Graham's anecdote on the Future Knowledge podcast, of a senior researcher at the New York Times thanking him personally for the Wayback Machine because there is material the researcher could not find in the Times's own archives, is a useful reminder of how dependent even the most well-resourced newsrooms are on the institution their corporate parents are simultaneously blocking (Freeland, 2026).

Academic researchers, particularly in digital humanities, computational social science, and information science, are structurally dependent on bulk access to archived collections. Zannettou et al. (2018), studying the use of web archiving services on social media, required access to millions of archived URLs to trace how archiving services are used and misused. Research of this kind is indistinguishable from AI training at the technical level; both involve programmatic bulk extraction. A policy that blocked all bulk extraction categorically would eliminate this research. The Credentialed Research Tier in Section 3 is designed to preserve it. As Kendra Albert observed on the podcast, researchers who study bias and discrimination on platforms are subject to the same body of scraping law that gets made by Reddit's lawsuits, which is precisely why the line must be drawn institutionally rather than technically (Freeland, 2026).

Commercial AI companies have a material interest in continued access to the Archive's collections, which represent an extraordinarily rich source of training data spanning three decades of the open web. There is also a counterintuitive argument for accommodating that interest. The Archive's collections are valuable to them precisely because the data is curated, deep, and long-running. A policy that closes the door entirely will push them toward worse data sources and will not stop the pressure they exert on the Archive's infrastructure. If commercial AI companies cannot access the Archive, they do not stop building models. They train on Reddit scrapes, low-quality social media data, and SEO-optimized junk content. A policy that opens a structured commercial channel, with attribution and reciprocity, can convert pressure into participation. A model trained on the Archive's curated historical collections, primary sources, and verified journalism may be more accurate and less biased than one built on whatever is easiest to extract from the open web. The structural parallel between AI training and academic text and data mining means that a network-layer block on commercial AI hits academic computation just as hard. The Data Trust Tier in Section 4 threads this needle by conditioning commercial access on transparency, provenance, and reciprocity rather than excluding it categorically.

Commercial scraping intermediaries like Automatio, Octoparse, ScrapeStorm, and the open-source Wayback Machine Scraper are a related but distinct group. Automatio openly markets "Intelligent Retry Logic" as a feature for working around the Archive's rate limits and IP blocks. The policy is not designed to accommodate this group. It is designed to make the technical workarounds these intermediaries provide legally costly to use against the Archive.

The proposed policy engages values articulated by Schafer and Winters in their framework of web archive governance: participation, consensus, accountability, transparency, effectiveness, inclusivity, and adherence to the rule of law (Schafer & Winters, 2021). The NASCIO and Council of State Archivists State Archiving Playbook (NASCIO & CoSA, 2018) establishes a parallel principle for electronic records management: records should be appropriately identified, properly classified, maintained according to retention policies, stored in a way that makes them locatable, and effectively preserved. The Archive's policy must meet those standards for its own collections before it can credibly require commercial users to do the same.

The Journal for Scholarly Communication, Librarianship, and Open Knowledge volume frames the librarians' professional commitment as access as praxis--access is not a value to declare but an active institutional commitment to maintain (Bonn, Bolick, & Cross, 2023). They advocate that "libraries as the ethical alternative" should shift from passive banks of resources to active stewards of ethical infrastructure, supporting the Archive's claim to a governance role in AI training data. This opinion is not reflected across the board, as stated in *Library and Archives 101: AI and the False Promise of Control*,

“Asserting control, demanding compensation, and conditioning access on the institution’s ability to dictate the terms of downstream use—that is not what libraries and archives are for.” (Hansen, 2026). However, it is recognized that archives will not “out-monopoly the monopolists by asserting control over their collections.” But rather by ensuring that the “institutions, researchers, and communities they serve have a seat at the table.” The Harvard Library Public Domain Corpus, released in late 2025, demonstrates the same logic from a different direction. By offering documented, legally sound public domain training data, Harvard shifted the power dynamic away from proprietary scraping and toward public stewardship. The Internet Archive is positioned to make a parallel claim, but only if it has the governance infrastructure to back it up.

Finally, ordinary users of the Wayback Machine, people who search for deleted web pages, archived news articles, and historical versions of public websites, are the Archive's largest constituency and its most diffuse. They have no formal organizational voice. Their interests are represented in the Archive's foundational mission. Any policy that imposes barriers on individual access to preserved web content would betray that mission. The Open Access Tier in Section 2 is specifically designed to ensure that the credentialing requirements imposed on bulk users do not, and cannot, spill over into restrictions on individual access.

IV. Obstacles and Challenges

A. Legal Uncertainty

The most significant legal obstacle is that the question of whether AI training on copyrighted archived content is fair use remains genuinely unsettled. More than one hundred active copyright cases are working through U.S. courts, and the outcome of any one of them could alter the landscape significantly. The Archive's preservation activity is on solid ground under *Authors Guild v. HathiTrust*, 755 F.3d 87 (2d Cir. 2014), and the broader text and data mining line including *Authors Guild v. Google*, 804 F.3d 202 (2d Cir. 2015), and *A.V. ex rel. Vanderhye v. iParadigms*, 562 F.3d 630 (4th Cir. 2009). The harder question is what happens when the Archive continues to operate an open commercial bulk-access channel after the 2023 incident. *Bartz v. Anthropic* and *Kadrey v. Meta* (both N.D. Cal. 2025) push the analysis in different directions, and *Thomson Reuters v. Ross*, 765 F. Supp. 3d 382 (D. Del. 2025), creates risk where AI outputs substitute for the source. The policy is therefore drafted at a level of principle in several places, with the expectation that implementation guidelines will be updated as the law develops.

The CMI removal question under Section 1202 is similarly unsettled. It is possible that a court will hold that CMI stripping during bulk extraction is independently actionable even where the underlying

use is fair use. The proposed policy's CMI preservation requirement in both the Credentialed Research Tier and the Data Trust Tier is designed to mitigate this risk prospectively. The Computer Fraud and Abuse Act (CFAA) also leaves a gap. *hiQ Labs v. LinkedIn* (9th Cir. 2022) and *X Corp. v. Bright Data* (N.D. Cal. 2024) hold that public data alone does not give rise to a CFAA claim, but the 2023 conduct, sustained, distributed, post-block, sits closer to a denial-of-service event than to ordinary scraping. As Albert observed on the Future Knowledge podcast, web scraping law is much less favorable than fair use law because it does not typically ask why the access is occurring. The policy must publish bulk-access conditions in the form courts can actually enforce, supported by registration under *ProCD v. Zeidenberg*, 86 F.3d 1447 (7th Cir. 1996), so that violations have a contractual basis.

B. Technical Limitations

The proposed policy's enforcement depends on technical capabilities the Archive does not currently have at full scale. Rate limiting, bot detection, and IP blocking are deployable with existing infrastructure, but distinguishing a credentialed academic researcher's bulk extraction from an uncredentialed AI company's bulk extraction requires credential verification systems, API key management, and real-time monitoring capabilities that require investment to build and maintain.

The ai.txt standard is promising but not yet mature. There is currently no universal requirement for AI companies to respect machine-readable signals, and voluntary adoption by bad-faith actors is not realistic. The Archive can implement ai.txt as part of its signaling infrastructure while acknowledging that it will not by itself stop determined bad actors. NC State University Libraries and similar institutions have demonstrated that network-level tiering, in which institutional users receive priority access and unidentified automated bots can be rate-limited or blocked, is technically achievable for organizations with their own network infrastructure. Adapting this model for the Archive's open internet-accessible architecture requires different engineering solutions, including content delivery network controls and API gateway management.

Lerman et al.'s Large-scale online deanonymization with LLMs (arXiv, February 2026) compounds the technical challenge on the data side, as even content the Archive treats as public can be cross-referenced with auxiliary data and large language models to produce identifications the original publication never contained. Hansen (2026) makes the larger point bluntly: technical control alone is a false promise. Policy must do most of the work that technology cannot. Technical access controls support enforcement; they do not replace governance.

C. Institutional Capacity

The Internet Archive is a nonprofit with limited administrative staff. Implementing a credentialing program requires staff time to review applications, negotiate agreements, conduct audits, and manage compliance. The proposed thirty-business-day review window for Credentialed Research Access is achievable for a small number of applications but may not scale if demand is substantial. The Archive should consider a consortium model, in which multiple memory institutions share credentialing infrastructure, to reduce per-institution costs. The ICPSR model, used by the Inter-university Consortium for Political and Social Research and described by Locke and Weber in the *Scholarly Communication, Librarianship, and Open Knowledge* volume, provides a precedent for shared academic data governance at scale (Bonn, Bolick, & Cross, 2023).

V. Decision Points and Proposed Resolutions

Decision 1. Scholarly and commercial AI use are technically the same. Should the policy treat them differently? Yes, but not at the network layer. The technologies overlap; the institutional purposes do not. The proposed resolution follows the UVA Archival AI Protocol's distinction between retrieval-based AI, which keeps source materials under institutional control, and training-based AI, which absorbs materials into model weights that cannot be reversed. Retrieval is presumptively allowed for both scholars and commercial parties under appropriate terms. Training is presumptively restricted unless provenance, attribution, and a real right to stop further use are guaranteed. This protects the academic researcher and the commercial AI lab differently because their accountability structures differ, not because their methods do. The trade-off is real: drawing the line institutionally rather than technically requires registration, documentation, and audit machinery the Archive does not currently operate. A purely technical line was rejected because the technology cannot draw it.

Decision 2. Should the Archive Implement Any Restrictions on Bulk Access? Yes. The case for unrestricted bulk access is no longer defensible. Documented evidence shows that uncontrolled extraction has crashed the Archive's infrastructure, triggered publisher withdrawals that threaten the preservation mission, and enabled commercial exploitation without consent or reciprocity. The narrow argument that any restriction undermines open access principles proves too much; it would equally prohibit rate limiting, a measure every major web service implements as a matter of infrastructure survival.

Decision 3. How should scholarly access be distinguished from commercial access? Through a Researcher Passport modeled on the ICPSR restricted-data access framework Locke and Weber describe in Bonn, Bolick, and Cross's *Scholarly Communication Librarianship and Open Knowledge*. The Passport

requires verifiable institutional affiliation or comparable credentials, a brief data management plan, attribution commitments, and rate limits calibrated to legitimate research volumes. It is not gatekeeping by topic; researchers are not turned away for what they want to study. It is gatekeeping by accountability, and it produces the documented record that strengthens the user's own fair use defense under HathiTrust and iParadigms. The Passport explicitly includes non-university pathways for independent researchers, journalists, local historians, community archivists, and students without institutional backing, so that the credentialing requirement does not function as exclusion. I am still a bit uncertain about this for independent researchers or journalists who desire to stay anonymous.

Decision 4. How Should the Archive Respond to Publisher Blocking? The Archive should proactively approach publishers who have implemented blocks and offer to implement publisher-requested AI training restrictions within its bulk access agreements. This is not a concession; it is the Archive's strongest argument for reversing the blocking trend. If the Archive can credibly demonstrate that bulk access to archived versions of publisher content is governed by agreements that prohibit AI training without publisher consent, the publisher's primary objection to archival crawling is addressed. Some publishers may not be persuaded, and some may maintain blocks for strategic reasons unrelated to AI. But the Archive should not accept the current trajectory as permanent. The Pointer Institute's Today's News for Tomorrow project, which Graham described on the Future Knowledge podcast as serving more than three hundred local newsrooms, is evidence that constructive partnership at scale is possible (Freeland, 2026).

Decision 5. Should the Archive Require Reciprocity from Commercial Users? Yes. The reciprocity requirement is the most forward-looking element of the policy and the element most likely to be contested by commercial AI companies. It is also the element best supported by the institutional logic of the Archive's mission. Brewster Kahle's anti-hallucination proposal, under which AI companies implement citation features that link chatbot outputs back to archived sources, is both a form of reciprocity and a direct contribution to the Archive's mission of providing access to a verifiable historical record. Reciprocity does not have to mean paying a fee. The Harvard Library Public Domain Corpus model demonstrates that public institutions can condition AI access on documented, legally sound terms without abandoning their open access principles. Tarkowski (2025) frames reciprocity and compensation as one of six discrete focus areas for AI data governance, and the Library Copyright Alliance's principles, while skeptical of compulsory licensing, do not foreclose voluntary commercial agreements (Library Copyright Alliance, 2023). The Archive should follow these precedents.

Decision 6. How Should Privacy Obligations Be Handled? Where archived content includes personal data subject to the California Consumer Privacy Act, comparable state privacy laws, or the General Data

Protection Regulation in jurisdictions where it applies, the Archive will maintain a workable process for handling valid deletion requests. Data Trust counterparties must implement parallel deletion processing within ninety days of receiving notice from the Archive that an upstream deletion request affects content within their licensed corpus. The CARE Principles for Indigenous Data Governance, the ICPSR framework, and the NASCIO-CoSA State Archiving Playbook all support treating de-identification as a documented process subject to ongoing review rather than as a one-time technical step (NASCIO & CoSA, 2018). Lerman et al. (2026) makes this a practical operational obligation, not a theoretical concern.

VI. Annotated Path Forward

In this part I will provide annotations for each section in the Proposed Policy Text: Automated and Bulk Access to Internet Archive Collections.

Section 1. Purpose and Scope.

***Annotation.** The 10-requests-per-second threshold is the central operational decision in the policy. It is calibrated to sit comfortably above the rates of standard scholarly text-mining workflows and well below the burst rates of the May 2023 incident, which reached tens of thousands of requests per second. The carve-out for unauthenticated browsing is the most important provision in the entire Policy. It ensures the Archive's mission is not implicated by the bulk-access governance framework. Defining every term that bears legal weight in the document itself, rather than leaving them to later interpretation, follows the U.S. Copyright Office's drafting practice in its Copyright and Artificial Intelligence, Part 3 report (May 2025).*

Section 2. Open Access Tier.

***Annotation.** The Open Access Tier is what most users will ever encounter. The credentialing requirements imposed on bulk users do not, and must not, spill over into restrictions on individual access.*

Section 3. Credentialed Research Tier (Academic and Public-Interest Bulk Access).

***Annotation.** The Credentialed Research Tier is adapted from the ICPSR restricted-data access framework that Brandon Locke and Nic Weber describe in Maria Bonn, Josh Bolick, and Will Cross, eds., Scholarly Communication Librarianship and Open Knowledge (ACRL 2023). The non-university pathway responds directly to the risk that credentialing becomes exclusion. The CMI preservation requirement is the policy's response to the Section 1202 issues raised in Kadrey v. Meta and related cases. The data-management-plan requirement produces the documented record that strengthens the user's own fair use defense under Authors Guild v. HathiTrust and A.V. ex rel. Vanderhye v. iParadigms.*

Section 4. Data Trust Tier (Commercial and Foundation Model Training).

Annotation. The four reciprocity options are deliberately structured to be flexible. Brewster Kahle's publicly proposed anti-hallucination citation linkage, in which AI outputs cite Archive sources and route users to them, is the most elegant of the four because it converts a commercial AI company from an extractor of archival value into a distributor of archival access. Alek Tarkowski's Data Governance in Open Source AI (Open Future and OSI, January 2025) names reciprocity and compensation as one of six discrete focus areas for AI data governance, and the Library Copyright Alliance's 2023 principles, while skeptical of compulsory licensing, do not foreclose voluntary commercial agreements. The market-substitute provision in (f) is the policy's response to the fourth fair use factor as it has developed in *Thomson Reuters v. Ross*. It does not protect publishers from AI competition in the abstract; it protects the Archive from being held responsible for that competition. The calibration of the right-to-stop in (d) follows the UVA Archival AI Protocol's distinction between narrow models, where decommissioning is realistic, and large mixed-corpus foundation models, where it is not. Demanding decommissioning of a frontier foundation model would discredit the policy. The narrower rights are credible and likely to be agreed to by a serious commercial counterparty.

Section 5. Machine-Readable Signaling.

Annotation. Machine-readable preference signals are incorporated as contractual conditions, not as standalone protections. Creative Commons' AI training preference signals work and the W3C TDM efforts are still maturing. The European Union's copyright directive already recognizes machine-readable TDM opt-outs as legally significant; U.S. law does not yet, and ai.txt as a voluntary standard carries no legal weight in U.S. courts. Embedding compliance with such signals into the Data Trust Agreement, with verification in the annual compliance report, converts a voluntary norm into a binding term. This is the bridge the Archive can offer publishers.

Section 6. Rate Limits and Infrastructure Protection.

Annotation. Designating these measures as TPMs under § 1201 strengthens the Archive's position against bad-faith circumvention, including the kind of automated retry logic that commercial scraping vendors openly market.

Section 7. Publisher-Requested Restrictions.

Annotation. Section 7 is the relational mechanism by which the legal and technical provisions translate into restored publisher trust. The Guardian's existing arrangement, in which the Archive limits access to some Guardian journalism while leaving regional homepages and topic pages open, is one model baseline. The Pointer Institute's Today's News for Tomorrow project, which Mark Graham described on

the Future Knowledge podcast as serving more than three hundred local newsrooms, demonstrates the same logic at the local-newsroom scale. The Archive cannot compel access, it can only offer credibility.

VII. Sources Cited

Allan, Darren. "News sites are blocking the Internet Archive's Wayback Machine." *TechRadar*, April 13, 2026.

Authors Guild v. HathiTrust, 755 F.3d 87 (2d Cir. 2014).

Authors Guild v. Google, Inc., 804 F.3d 202 (2d Cir. 2015).

Automatio. *How to Scrape Archive.org / Internet Archive Web Scraper*. Vendor documentation.

A.V. ex rel. Vanderhye v. iParadigms, L.L.C., 562 F.3d 630 (4th Cir. 2009).

Bartz v. Anthropic PBC, No. 3:24-cv-05417-WHA (N.D. Cal. 2025).

Bonn, Maria, Josh Bolick, and Will Cross, eds. *Scholarly Communication Librarianship and Open Knowledge*. Chicago: Association of College and Research Libraries, 2023. Including Brandon Locke and Nic Weber, "Ethics of Open Data."

California Consumer Privacy Act, Cal. Civ. Code §§ 1798.100 et seq.

Communications Decency Act, 47 U.S.C. § 230, including § 230(e)(2).

Computer Fraud and Abuse Act, 18 U.S.C. § 1030.

Copyright Act of 1976, 17 U.S.C. §§ 101 et seq., including §§ 106, 107, 108, 1201, 1202.

Deck, Andrew. "Why news sites are blocking the Wayback Machine." *Marketplace / Nieman Lab*, April 21, 2026.

Doe v. GitHub, Inc. (9th Cir., on appeal).

Electronic Frontier Foundation, Fight for the Future, and Public Knowledge. Open journalist letter of support to the Internet Archive, April 2026.

Freeland, Chris. "Wayback Machine Director: We Are 'Collateral Damage' in the Fight Between AI Companies and Publishers." *Internet Archive Blog*, May 6, 2026.

<https://blog.archive.org/2026/05/06/wayback-machine-director-we-are-collateral-damage-in-the-fight-between-ai-companies-and-publishers/>

Hachette Book Group, Inc. v. Internet Archive (2d Cir. 2024).

Hansen, Dave. *Library and Archives 101: AI and the False Promise of Control*. Authors Alliance, March 16, 2026.

hiQ Labs, Inc. v. LinkedIn Corp., 31 F.4th 1180 (9th Cir. 2022).

Internet Archive. Scraping API documentation. <https://archive.org/services/search/v1/scrape>

Joseph, Jibin. "Reddit Is Blocking the Internet Archive's Wayback Machine." *PCMag*, August 12, 2025.

Kadrey v. Meta Platforms, Inc., 2025 WL 1786418 (N.D. Cal. June 27, 2025).

Kahle, Brewster. "Let us serve you, but don't bring us down." *Internet Archive Blog*, May 29, 2023.
blog.archive.org.

Knibbs, Kate. "The Internet's Most Powerful Archiving Tool Is in Peril." *Wired*, April 13, 2026.

Kennedy, Dan. "News publishers are blocking the Wayback Machine." *Media Nation*, April 2026.

Lemley, Mark A., and Peter Henderson. *The Mirage of Artificial Intelligence Terms of Use Restrictions*. 2024.

Lerman, Simon, Davide Paleka, John Swanson, Manuel Aerni, Nicholas Carlini, and Florian Tramèr. "Large-scale online deanonymization with LLMs." *arXiv preprint arXiv:2602.16800v2*, February 25, 2026.

Library Copyright Alliance. *Principles for Copyright and Artificial Intelligence*. July 10, 2023.

Lo, Leo S. *The University of Virginia Archival AI Protocol*, v.1.1, with companion Adoption Guide and Implementation Toolkit. University of Virginia Library, January 27, 2026.
<https://doi.org/10.18130/5dqf-9w86>

Lovejoy, Ben. "News sites blocking the Wayback Machine." *9to5Mac*, April 13, 2026.

Messarra, Luca, Chris Freeland, and Juliya Ziskina. *Vanishing Culture: A Report on Our Fragile Cultural Record*. San Francisco: Internet Archive Press, 2026.

Moody v. NetChoice, LLC, 603 U.S. 707 (2024).

NASCIO and Council of State Archivists. *State Archiving Playbook*. State Archiving in the Digital Era: A Playbook for the Preservation of Electronic Records. NASCIO-CoSA Pub. 18, October 2018.

Originality AI. Analysis of Internet Archive crawler blocking by major news publishers. January 2026.

Pew Research Center. "When Online Content Disappears." May 2024.

"Preserving the Web in the Age of AI." *Future Knowledge* (podcast), with Mark Graham, Mike Masnick, Kendra Albert, hosted by Chris Freeland and Dave Hansen. Internet Archive and Authors Alliance, May 6, 2026.
<https://futureknowledge.transistor.fm/episodes/preserving-the-web-in-the-age-of-ai>

ProCD, Inc. v. Zeidenberg, 86 F.3d 1447 (7th Cir. 1996).

Radauskas, Gintaras. "NYT, USA Today among major news sites blocking the Internet Archive's Wayback Machine: Why?" *Cybernews*, April 15, 2026.

Richard, Isaiah. "Internet Archive Hit by Cyberattack." *Tech Times*, October 14, 2024.

Rooney, John. "Best Internet Archive Scraper 2024." *WebScrapingSite*.

Sag, Matthew. "Copyright and Copy-Reliant Technology." *Northwestern University Law Review* 103 (2009): 1607.

Sangaline. *Wayback Machine Scraper*. GitHub repository.

Schafer, Valerie, and Jane Winters. "The values of web archives." *International Journal of Digital Humanities* 2 (2021): 129-144.

Singla, Anmol. "Why news sites are blocking the Internet Archive." *Firstpost*, April 14, 2026.

Sircar, Anisha. 'Why Major News Sites Are Blocking the Internet Archive's Wayback Machine.' Forbes/Yahoo Finance, April 14, 2026.

Society of American Archivists. *SAA Core Values Statement and Code of Ethics*. February 2005, last revised August 2025.

Tarkowski, Alek. *Data Governance in Open Source AI: Enabling Responsible and Systemic Access*. Open Future and Open Source Initiative, January 22, 2025.

Thaler v. Perlmutter, No. 23-5233 (D.C. Cir. 2025).

Thomson Reuters Enterprise Centre GmbH v. Ross Intelligence Inc., 765 F. Supp. 3d 382 (D. Del. 2025).

University of Virginia. UVA Archival AI Protocol (AAIP). 2024-25.

U.S. Copyright Office. *Copyright and Artificial Intelligence, Part 3: Generative AI Training* (Pre-Publication Version). May 2025.

U.S. Copyright Office. *Section 1201 Rulemaking: Recommendation of the Register of Copyrights*. 2024.

Vienna Principles: A Vision for Scholarly Communication. 2016.

Washington Post analysis of Google's C4 dataset, 2023.

X Corp. v. Bright Data Ltd., No. C 23-03301 JSW (N.D. Cal. Apr. 1, 2024).

Zannettou, Savvas, Jeremy Blackburn, Emiliano De Cristofaro, Michael Sirivianos, and Gianluca Stringhini. "Understanding Web Archiving Services and Their (Mis)Use on Social Media." *Proceedings of the International AAAI Conference on Web and Social Media (ICWSM)*, 2018.