

Authorship and Creativity in the Age of Generative AI

75 minutes total

Summary:

This facilitation introduces the class to questions about authorship and creativity in a time when generative AI can produce convincing images and text. We propose that authorship in the age of generative AI cannot be reduced to output ownership alone; it must be understood as a contested relationship among human intention, machine regulation, training-data extraction, and legal recognition. We start by talking through what creativity means and why intention and judgment matter for human creators. Then we give a quick overview of how AI systems are built on large datasets, human labor, and unequal access to resources. After that, students work in small groups with two case studies that show very different experiences of AI in creative work: one involving a major corporation and one involving independent artists. The goal is for students to compare these situations and think about who is doing the creative work, who gets credit or compensation, and what authorship should look like when AI is involved.

Before Class Readings:

- Boden, M. A. (2014). Creativity and artificial intelligence. [Link](#)
- Moulaison-Sandy, H. (2023). What is a person? Emerging interpretations of AI authorship and attribution. *Proceedings of the Association for Information Science and Technology*, 60(1), 279–290. [Link](#)
- Optional read: Thaler v. Perlmutter [Link](#)

<https://law.justia.com/cases/federal/appellate-courts/cadc/23-5233/23-5233-2025-03-18.html>

Question: What does authorship mean in a world of generative AI?

Intro and framing for our topic (~10 minutes):

Led by: Meagan

Content (Key Points):

- Briefly introduce the topic: authorship and creativity in the age of generative AI.
 - Start by asking the class what comes to mind when they hear the word creativity. Let a few students respond. Then move from those intuitive answers into a concise working definition.
 - Creativity involves more than novelty; it includes intention, judgment, and context.
 - Humans make choices about meaning and value. AI does not.
- Explain why generative AI complicates traditional ideas of creativity:
 - AI can produce outputs that look creative, but AI is remixing patterns from training data, not forming intentions.
- Raise the central authorship question:

- When AI generates something, who is the “author”?
 - The prompter, the company, the original creators whose work trained the model, or nobody in the traditional sense.

Ideas for light engagement (Maybe prompting them to turn to a neighbor or 2 instead of moving into small groups):

- Ask the class what comes to mind when they think of “creativity.”
- Ask whether AI feels creative to them and why.
- Questions for thought:
 - When a novelist uses a thesaurus to find exactly the right word, do they own that word choice?
 - When a photographer uses Photoshop to composite two images together, do they own the result?
 - When a graphic designer types a detailed prompt into Midjourney and generates an image, do they own that image?

Evidence for this section:

- Boden’s definition of creativity (novel, surprising, valuable). [Link](#)
 - What is Creativity? Boden uses a three-part framework to answer this question. He says a creative work is novel, surprising, and valuable.
 - Creativity in humans, in particular, involves intention. A human chooses what to make, why it matters, and what it means.
 - Contrast this with generative AI, which produces outputs by statistically predicting plausible continuations of training data. AI can achieve novelty and surprise; whether it can achieve value in the full sense (including social, ethical, and intentional value) becomes the core question.
- Moulaison-Sandy on why institutions reject AI authorship. [Link](#)

Purpose: Establish topic and lay groundwork for understanding and analyzing the case studies later on.

Mini lecture/background for case studies: AI as an extractive system (~10 minutes)

Led by: Max

Content (Key Points):

- Explain how generative AI systems are built: *Hrishika*
 - Models are trained on massive datasets of human-created material (images, writing/text, music, code, etc.).
 - They learn patterns, associations, and stylistic signatures at a statistical level. When prompted, they remix those patterns to generate new outputs. It does not retrieve stored content, understand what it is making, or have any stake in the outcome.

- Much of this data is scraped from the internet without consent.
- Highlight the invisible labor behind AI (starting to take a stance here leaning into the argument that AI should not be considered an author):
 - Data labeling, content moderation, dataset cleaning. All often low-paid or unpaid.
- Emphasize that AI systems are not neutral:
 - They reflect the power structures and inequalities of the world they're built in.
 - Large corporations can negotiate licenses while independent creators usually cannot.
- Connect this to fairness and ownership:
 - If AI is built on human creative labor, what obligations do companies have to those creators?
 - Who benefits from AI systems, and who is left out?

Evidence for this section:

- Pasquinelli & Joler (2021): Nooscape Manifested (AI as instrument of knowledge extraction). [Link](#)
 - AI absorbs and compresses human-generated creative work for commercial exploitation without crediting, compensating, or notifying the people whose work made the system possible.
- Terranova (2000) on free labor in digital economies. [Link](#)
 - Digital platforms have always run on "free labor": the unwaged creative and affective work users contribute, which generates value for others without compensation.
 - Gen AI is only new due to its scale and directness. Rather than it just being your direct contribution (ie. a post or comment) that makes a platform valuable, it is now that your artistic style, compositional choices, and years of craft are absorbed into a system that then competes with you commercially.
- Birhane (2021) on algorithmic injustice and relational ethics. [Link](#)
 - The infrastructure of AI creativity is built partly on the uncompensated labor of the people least positioned to benefit.

Purpose: Provide more background on the ethical and economic contexts before moving into the case studies.

Brief Legal Framing: How the courts are currently answering the authorship question (~10 minutes)

Led by: Hrishika

Content (Key Points):

- Explain why it is important to answer this question about authorship.

- “Why does any of this matter beyond philosophy?”
Authorship determines ownership. And ownership determines who gets paid, who gets credit, and who has legal recourse when something goes wrong.
- “Who cares?”
This matters now because generative AI is no longer just a futuristic issue. It is already reshaping how people write, design, illustrate, research, and publish. And legally, institutions are trying to respond in real time.

So Max just walked us through how AI systems are built on extracted human labor and creative work. That raises the obvious next question: does any of that actually matter legally? And the answer is yes — it matters a lot. Because authorship isn't just a philosophical concept. It's the mechanism that determines who gets paid, who gets credit, and who has legal recourse when something goes wrong.

And this is no longer a futuristic problem. Generative AI is already reshaping how people write, design, illustrate, research, and publish. Courts and institutions are trying to respond in real time — and they don't always agree on the answers. So let's look at how the law is currently thinking about this.

- How are institutions currently answering the question?
 - Explain: We are not spending long on copyright, but it helps to see how the law is currently answering the question we are asking.
 - The Purpose of Copyright: “[C]opyright law ultimately serves the purpose of **enriching the general public through access to creative works**. . . . The statute achieves that end by striking a balance between two subsidiary aims: encouraging and rewarding authors' creations while also enabling others to build on that work.” (Kirtsaeng v. John Wiley & Sons, 136 S. CT. 1979, 1986 (2016))
 - Creative works often require: Time (months or years of labor), Skill (training, expertise, creative judgment), Risk (no guarantee of success), Upfront investment (equipment, research, collaborators).
 - We want to encourage more creative works. In a purely free market, most people will simply copy other creative work and recreate it cheaply. Copyright de-incentivizes this by protecting a person's investment of creative labor.
 - However, AI flips the entire incentive structure. We're not talking about Photoshop or AutoTune, tools that assist human creativity but leave the human clearly in control. Generative AI systems can receive a short text prompt and produce a fully rendered image, a complete musical composition, or a polished essay. The human provides what might be called "creative direction." The system does the rendering.

- If AI outputs are: cheap to generate, unlimited, not tied to human effort. Then granting exclusivity to each output: doesn't incentivize creation, may clog cultural space, may create monopolies over abundance.

Copyright law was designed with a specific purpose: to enrich the public by encouraging creative work. The way it does that is by striking a balance — it rewards authors for their investment of time, skill, and risk, while also making sure others can eventually build on that work.

Think about what it actually takes to make something creative. Months or years of labor. Specialized skill and training. The risk that no one will care about what you made. Often real upfront costs — equipment, research, collaborators. Copyright says: we want to encourage more of this, so we'll protect your investment.

But generative AI disrupts that entire logic. We're not talking about Photoshop or AutoTune — tools that assist human creativity while leaving the human clearly in control. A generative AI system can take a short text prompt and produce a fully rendered image, a complete musical composition, or a polished essay. The human provides creative direction. The system does the rendering. So if AI outputs are cheap, unlimited, and not tied to human effort — what exactly are we incentivizing by granting copyright protection?

- Introduce: *Thaler v. Perlmutter*.
 - Stephen Thaler tried to register an artwork that he said was created autonomously by his AI system, and the courts rejected the claim because copyright law still requires human authorship.
- How/why does the law explicitly require a human author?
 - The human authorship requirement goes back to the constitutional basis of copyright itself, which was designed to promote "the progress of science and useful arts", and is a goal that has always been understood to mean incentivizing human creativity.
 - They also posit that machines have no lifespan, dependents, etc. that would otherwise be considered for an author to receive copyright. There exists no basis for the economic protections copyright was designed to provide.

The most directly relevant case we have right now is *Thaler v. Perlmutter*. Stephen Thaler built an AI system he called the "Creativity Machine," which autonomously generated a two-dimensional image. He submitted it for copyright registration listing the AI as the sole author. The Copyright Office rejected it. He sued. And the courts — first the district court in 2023, then the D.C. Circuit in 2025 — agreed with the Copyright Office.

The reasoning matters. Copyright traces back to the constitutional goal of promoting the progress of the arts and sciences — and that goal has always been understood as incentivizing *human* creativity. The courts also pointed out something practical: copyright's protections were designed around human economic interests. Authors have livelihoods, dependents, finite careers. A machine has none of those things. There's no basis for the protections copyright was designed to provide.

So the law has answered one question clearly: AI cannot be an author. But it left a much harder question completely open — and that's the one that actually affects the people in our case studies. How much human creative input is enough to qualify for protection in an AI-assisted work? And what about the copyrighted material that went *into* training these systems in the first place? That's still being litigated. Which brings us to fair use.

- Introduce the 4 statutes of Fair Use that determine and assign authorship to works derived from other works.
 - 1. The purpose and character of the use. Is the use commercial or educational? Is it transformative?
 - Courts look favorably on uses that genuinely transform the original. AI companies argue that training is transformative because the model learns patterns rather than reproducing content. Critics argue that commercially training on art to build a commercial image generator is not transformative, but merely reproduction at scale.
 - 2. The nature of the copyrighted work. Is the original work highly creative, like an illustration or painting, or more factual, like a database?
 - More creative works get stronger protection. This factor tends to favor artists whose distinctive, original styles were scraped to train image generators.
 - 3. The amount of the work used. How much of the original was copied, a snippet or the whole thing?
 - AI training typically uses entire images or entire texts. That tends to cut against a fair use finding, though AI companies argue no single work is reproduced in the output.
 - 4. The market effect. Does the use harm the market for the original?
 - If an AI system can generate images "in the style of" a specific living artist and those generated images compete with and undercut that artist's livelihood, the courts would find market harm.
 - No single factor is decisive. Courts weigh all four together. And as of right now, we don't have a definitive ruling on whether training AI systems on copyrighted work is or isn't fair use. We will see a case study soon that delves into this issue.
- Takeaway: The courts said AI cannot be an author. Did not explicitly define how much human authorship is required in AI-assisted creative works, except that it must be significant.
- Four numbered boxes:

- **Purpose & character** — Is it transformative? Commercial use weighs against fair use
 - **Nature of the work** — More creative originals get stronger protection
 - **Amount used** — AI typically ingests entire works; cuts against fair use
 - **Market effect** — Can AI outputs compete with and undercut the original creator's livelihood?
- Bottom note: *No single factor is decisive. No definitive ruling yet on AI training data.*

To understand the case studies you're about to work through — especially *Andersen v. Stability AI* — you need to know about fair use. Fair use is the legal doctrine that allows limited use of copyrighted material without permission, under certain conditions. Courts weigh four factors together to decide.

First: the purpose and character of the use. Is it transformative — does it do something genuinely new with the material? AI companies argue that training transforms data into statistical patterns, not reproductions. Critics argue that using artists' work to build a commercial image generator isn't transformation — it's reproduction at scale.

Second: the nature of the copyrighted work. More creative, expressive works get stronger protection. This tends to favor artists whose distinctive styles were scraped without consent.

Third: how much of the work was used. AI training typically ingests entire images and texts. That generally cuts against a fair use finding, though companies argue no single work appears in the output.

Fourth, and often most important: market effect. If an AI system can generate images "in the style of" a specific living artist, and those images compete with and undercut that artist's market — courts take that seriously.

No single factor is decisive, and they're always weighed together. Crucially, as of right now, there is no definitive court ruling on whether training AI systems on copyrighted work constitutes fair use. That's exactly what's being fought out in the *Andersen* case. Keep these four factors in mind when you're in your small groups — they'll give you a legal vocabulary for what's at stake.

Purpose: Ground the philosophical and ethical discussion in live legal stakes. Introduce *Thaler v. Perlmutter* as a concrete anchor for the human authorship requirement, and walk through the four fair use factors efficiently so students have the vocabulary they need for the *Andersen* case study, to compare how their ethical frameworks compare against what is held in real life.

Case Studies: Small groups (~30 minutes)

Led by: Meagan

Content:

Divide class into groups. Have the class read case summaries provided by us and give them guiding questions to discuss in their groups. The class will have 10 minutes to read the case studies and 20 minutes to discuss guiding questions regarding the cases. We will notify the groups when the 10 minutes for the first case has passed to remind groups to switch if they have not already.

Case 1: Disney x OpenAI

Case 2: Andersen v. Stability AI

Case summary drafts:

1. Disney x OpenAI:

In 2024, Disney entered a partnership with OpenAI to explore the use of generative AI across its vast intellectual property ecosystem. The partnership has just recently ended, but we can still analyze its ethics. Unlike many AI controversies centered on unlicensed data scraping, this collaboration was built on explicitly negotiated terms: Disney grants OpenAI controlled access to its archives, characters, and creative assets, while retaining strong oversight over how its IP is used. The partnership aimed to accelerate content production, assist animators and writers, and develop new forms of interactive storytelling.

Disney's long history of aggressively defending its copyrights made this partnership especially notable. For decades, the company has lobbied for extended copyright terms and fought unauthorized uses of its characters. With the partnership, it was embracing generative AI as a tool for creative production, but only under conditions that preserve corporate control. This raised questions about how AI may reinforce existing power structures in creative industries, where large corporations can negotiate favorable terms while independent creators cannot.

The partnership also signaled a shift in how creativity is conceptualized, from a human-driven process to a hybrid model where AI assists or even generates content. Disney framed this as innovation, but critics argued it risks reducing human creativity to a managerial input rather than a central artistic force.

When thinking about the guiding questions below, imagine you are a creative worker inside Disney whose job may have been augmented or automated by AI.

Source: [Link](#)

2. Andersen v. Stability AI:

In 2023, a group of visual artists, including Sarah Andersen, filed a class-action lawsuit against Stability AI, Midjourney, and DeviantArt. The artists allege that their copyrighted works were scraped from the internet without consent and used to train image-generation models such as Stable Diffusion. These models can produce images that mimic the artists' distinctive styles, raising concerns about style appropriation, loss of livelihood, and uncompensated creative labor.

The plaintiffs argue that training on their work constitutes copyright infringement and that the resulting models function as tools for producing derivative works that compete with their own. Stability AI counters that training data is "fair use," and that the model does not store or reproduce specific images but learns statistical patterns. The case is ongoing, and courts have not yet established clear precedent for whether training data constitutes infringement.

Beyond the legal questions, the case highlights the vulnerability of independent creators in an AI-driven creative economy. Unlike Disney, individual artists cannot negotiate licensing agreements or control how their work is used. Their labor becomes part of a massive dataset that fuels commercial AI systems without credit, compensation, or consent.

When thinking about the guiding questions below, imagine you are an independent artist whose work was scraped without your knowledge.

Sources: [Link](#), [Link](#)

Guiding Questions may be:

1. Who is doing the creative work here, and who is getting the credit?
2. What feels fair or unfair about how human creators' work is being used?
3. Who benefits most, and who is most vulnerable?
4. What would a more ethical or balanced version of this scenario look like?

After discussing both case studies (if groups have time):

What similarities and differences do you notice between the two cases in terms of authorship, fairness, and power?

Whole class discussion of case studies (~10 minutes)

Led by: Max and Hrishika

- Start with the similarities and differences of the case studies.
- Based on the cases, what does authorship mean in a generative era?
- Can AI-generated work be considered creative? (Calling back to earlier content)
- What obligations do developers and corporations have to human creators?
- How do power and labor shape who gets credit and compensation?

Conclusion (~5 minutes)

Led by: Meagan and Hrishika

Content:

- Our take on AI authorship:
 - Based on the evidence we explored, our group believes that AI cannot meaningfully be considered an author in its current form. While generative models can produce outputs that may look creative, they do not have intention, judgment, or accountability, and their work depends entirely on human-created data and labor. Because of this, authorship today should remain tied to human responsibility and human decision-making.
- A suggestion for a framework moving forward ...
 - Without delving too far into complex legal matter, we believe a more useful framework for thinking about AI authorship would focus on transparency about how AI systems are built, clear attribution of the human roles involved (from dataset creators to prompt writers), and stronger protections for the people whose work and labor make these systems possible.

We will ask the class if there are any questions.



Authorship & **Cre**Ativity

in the age of Generative AI

What does authorship mean in a world of generative AI?

Today's Outline:

01

Initial Discussion

02

Intro & Framing

03

AI as an Extractive System

04

Brief Legal Framing


05

Conclusion

Which One Is Human made? Does It Matter?



AI slop



Claude Monet's

... contd



- Now that you know, does it matter?

Follow up

If the AI image looks convincing, should it have value equal to Picasso's?

- Who owns the AI version?
 - The person who typed the prompt?
 - The AI company?
 - No one?
 - Monet's estate?

Intentional

Producing
ideas

Open Mind

Surprising

Novel

What do we think of when we hear

Creativity?

Original

Choice

Unique

Innovation

Valuable

Useful

What is Creativity?

Novel

Something new and unexpected, not a simple recombination of the familiar.

Surprising

It challenges or shifts our existing understanding or expectations.

Valuable

It matters. Socially, ethically, or aesthetically.

Boden (2014)

Human creativity adds intention, judgment, and meaning. Humans choose what to make, why it matters, and what it means.

How AI Complicates Creativity

Human Creators

- Forms intentions
- Makes meaning-based decisions
- Exercises aesthetic judgment
- Has accountability for the work

VS

Generative AI

- Remixes patterns from training data
- Statistical predictions
- No stake in the outcome
- Accountability?

AI can achieve novelty and surprise, but can it achieve value in the full sense?

When AI generates something, who is the
“author” ?

Prompter?

*Directed the creative
vision*

Company?

*Built and own the
model*

Original Creators?

*Whose work trained
the model*


Nobody?

*Not in the traditional
sense*

Is AI creative?

- Trained on massive datasets (text, images, code, etc.)
- Learns patterns, styles, associations
- Generates by recombining patterns

*AI generates
without intention*



The Lovelace 2.0 Test

We propose a test designed to challenge the premise that a computational system can originate a creative artifact. We believe that a certain subset of creative acts necessitates human-level intelligence, thus rendering both a test of creativity and also a test of intelligence.

The Lovelace 2.0 Test is as follows: artificial agent a passes the Lovelace Test if and only if:

- a creates an artifact o of type t ,
- o conforms to a set of constraints C where $c_i \in C$ is any criterion expressible in natural language,
- a human evaluator h , having chosen t and C , is satisfied that o is a valid instance of t and meets C , and
- a human referee r determines the combination of t and C to not be impossible.

The constraints set C makes the test Google-proof and resistant to Chinese Room arguments. An evaluator is allowed to impose as many constraints as he or she deems necessary to ensure that the system produces a novel and surprising artifact. For example: “create a story in which a boy falls in love with a girl, aliens abduct the boy, and the girl saves the world with the help of a talking cat.” While C does not necessarily need to be expressed in natural language, the set of possible constraints must be equivalent to the set of all concepts that can be expressed by a human mind. The ability to correctly respond to the given set of constraints C is a strong indicator of intelligence.

The evaluation of the test is simple: a human evaluator is allowed to choose t and C and determine whether the resultant artifact is an example of the given type and whether it satisfactorily meets all the constraints. Aesthetic valuations are not considered. We suggest that the judge be allowed to repeat the test any number of times with different t and C .

The human referee r is necessary to prevent the situation where the judge presents the agent with a combination of t and C that are impossible to meet even by humans. The referee should be an expert on t who can veto judge inputs based on his or her expert opinion on what is known about t .

The Lovelace Test – Creativity as the Measure of Machine Intelligence

In simpler terms:

- The system must create something **novel, unexpected, and valuable**.
- The creation must be **unexplainable as a direct byproduct of programming instructions**.
- The creator (the programmer) should be genuinely **surprised** by the system’s output.

Thus, while the Turing Test measures **deceptive imitation**, the Lovelace Test measures **autonomous creativity**.


<https://benarous-academy.com/the-lovelace-test-creativity-as-the-measure-of-machine-intelligence/>

<https://arxiv.org/pdf/1410.6142v1>

II. What AI Actually "Learns"

Statistical Weights, Not Meaning

Every token (a word, word-fragment, punctuation mark, etc.) is assigned a **numerical weight** encoding how frequently it appears alongside other tokens. The model builds a record of correlation, not comprehension.




Human beings represent English words with a sequence of letters, like C-A-T for **cat**.

Language models use a long list of numbers called a **word vector**. For example, here's one way to represent **cat** as a vector: →

[0.026275847107172012, 0.010517630726099014, -0.03723245859146118, 0.2634263336658478, 0.27159205079078674, -0.004058598540723324, 0.09411223232746124, 0.03763202950358391, 0.0677901953458786, 0.013615828938782215, -0.16626013815402985, -0.07512251287698746, 0.03231512010097504, 0.01070499885827303, -0.1998470425605774, -0.1406830996274948, -0.35442763566970825, -0.2826053202152252, 0.20517906546592712, -0.061711959540843964, 0.0553741529583931, -0.07947561889886856, 0.36058270931243896, -0.18250514566898346, 0.09193558245897293, 0.25107425451278687, 0.0032481239177286625, -0.2228851467370987, 0.12824469804763794, 0.1804560273885727, -0.22922773659229279, -0.15449629724025726, -0.4489767253398895, -0.0011336790630593896, 0.15314607322216034, -0.5254043936729431, -0.06919074803590775, 0.07075195014476776, 0.04077175632119179, 0.10582190752029419, 0.21858002245426178, -0.16373135149478912, -0.18926094472408295, -0.057782456278800964, 0.13000017404556274, -0.254604309797287, 0.1952560693025589, -0.5137519836425781, 0.3309118449687958, -0.4543110728263855, 0.110003761947155, 0.13061614334583282, 0.4761994779109955, -0.10132287442684174, -0.37049153447151184, 0.11823231726884842, 0.4599807858467102, 0.14639972150325775, 0.12805140018463135, 0.2130066454410553, 0.09224042296409607, -0.12506945431232452, -0.05348290130496025, 0.020665831863880157, 0.04905964434146881, 0.05350551754236221, 0.08142400532960892, 0.13870194554328918, 0.02489626035094261, -0.1033845990896225, 0.35924115777015686, 0.48756086826324463, 0.07979368418455124, -0.26932069659233093, -0.1587565690279007, -0.21406079828739166, 0.14890775084495544, 0.5726392865180969, 0.050989676266908646, -0.2300463318824768, -0.5037204623222351, 0.2688310146331787, 0.009941095486283302, 0.04406436160206795, -0.34434473514556885, -0.005059537943452597, -0.06580474972724915, 0.03520386666059494, 0.046381719410419464, 0.1729619950056076, 0.0627308040857315, 0.13112123310565948, 0.020175093784928322, 0.08860735595226288, 0.00029903074027970433, -0.007828752510249615, -0.025910617783665657, 0.029675563797354698, -0.053532760590314865, -0.011369606480002403, 0.19633322954177856, 0.07438189536333084, -0.06906607747077942, -0.537365198135376, 0.11993257701396912, 0.006955144461244345, -0.004988966975361109, -0.21147902309894562, -0.2904081612627716, 0.016421818304092, -0.36078450083732605, 0.19249385595321655, 0.46174725890159607, -0.20630729198411111111, 0.06679052114486694, 0.06833336502313614, -0.14867301285266876, 0.03586705029011111111, -0.13063858449459076, 0.11563169211149216, -0.02872380241751671, -0.25222587585441111111, -0.13110628724098206, 0.12749220430850983, 0.016964489594101906, -0.06233049556111111111, 0.10840493440628052, 0.15071143209934235, 0.21202605962753296, -0.27901750802911111111, -0.01712099090218544, 0.231328085064888, 0.4775627553462982, -0.23476198315620111111, 0.06952516734600067, -0.0007713703089393675, 0.2684878706932068, 0.21182623501111111111, 0.14450983703136444, -0.3424213230609894, -0.2570018768310547, -0.21082445979111111111, -0.20065836608409882, 0.009108797647058964, -0.19227707386016846, -0.06321350111111111111, 0.3057820200920105, 0.18081414699554443, -0.19904017448425293, -0.0208013448111111111111, 0.29011908173561096, -0.2552326023578644, -0.09783080965280533, -0.31859758411111111111, -0.09839189797639847, -0.09000033885240555, 0.03903864324092865, 0.07868181111111111111]



III. Generation Is Prediction, Not Expression



<https://www.understandingai.org/p/large-language-models-explained-with>

http://vectors.nlpl.eu/explore/embeddings/en/MOD_enwiki_upos_skipgram_300_2_2021/cat_NOUN/

IV. Why This Matters for Authorship

Human Author

Has a reason for making something

Makes choices based on meaning

Can be held accountable for the work

Draws on lived experience and judgement

vs

Generative AI

Has no goal or intention

Makes choices based on probability

Has no stake in the output

Draws on statistical patterns in data

AI doesn't generate from understanding. It generates from statistical proximity to what humans have already made.

02: AI as an Extractive System

I. Free Labor in the Digital Economy

Terranova (2000): Free Labor

Digital platforms have always run on "free labor" – unwaged creative and affective work users contribute, generating value for others without compensation.

Gen AI is new only in scale and directness. Now your artistic style, compositional choices, and years of craft are absorbed into a system that then competes with you commercially.

The Core Argument

If AI is built on human creative labor, what obligations do companies have to those creators? Who benefits from AI systems, and who is left out?

II. Where the Data Comes From

Data is often:

Scraped from the internet

Unlicensed

Unconsented

Human creativity → Training data

III. Hidden Labor + Extraction

Behind AI:

Data labeling

Content moderation

Dataset cleaning

Extracts:

Style

Knowledge

Creative work

AI depends on human labor

*Terranova (2000)
Pasquinelli, M., & Joler, V. (2021)*

IV. Power, Fairness, and Authorship

Corporations

- Control systems
- Negotiate licenses

Individuals

- No consent
- No compensation

Who gets to be an “author”?

So...? Same question



03. Brief Legal Framing

I. Why Authorship Matters

Who gets paid

Compensation flows to whoever
holds the rights

Who get credit

Attribution shapes professional
reputation and legacy

Who has recourse

Only rights-holders can sue for
infringement

II. How are the courts answering the authorship question

The Purpose of Copyright

“[C]opyright law ultimately serves the purpose of enriching the general public through access to creative works. . . . The statute achieves that end by striking a balance between two subsidiary aims: encouraging and rewarding authors' creations while also enabling others to build on that work.”

(Kirtsaeng v. John Wiley & Sons, 2016)

AI Flips the Incentive Structure

Unlike Photoshop or AutoTune, generative AI can produce a fully rendered image or complete essay from a short prompt. If outputs are cheap, unlimited, and untied to human effort...

what are we incentivizing?

A Recent Entry to Paradise



Fair Use: How Courts Weigh Training Data Claims

No definitive ruling yet on whether training AI on copyrighted work constitutes fair use. Courts weigh all four factors together.

1

Purpose & Character

Is it transformative? Commercial use weighs against fair use. AI companies claim training transforms data into patterns; critics call it reproduction at scale.

2

Nature of the Work

More creative, expressive works get stronger protection — favoring artists whose distinctive styles were scraped without consent.

3

Amount Used

AI training typically ingests entire images and texts, which cuts against fair use — though companies argue no single work appears in outputs.

4

Market Effect

If AI generates images "in the style of" a living artist, competing with their livelihood — courts take that seriously.

Real World Case Studies

Case 1: Disney x OpenAI

Corporate control & licensed IP

Case 2: Andersen v. Stability AI

Independent creators & data scraping

Think about these questions as we talk through the cases:

1. Who is doing the creative work here, and who is getting the credit?
2. What feels fair or unfair about how human creators' work is being used?
3. Who benefits most, and who is most vulnerable?
4. What would a more ethical or balanced version of this scenario look like?

Case 1: Disney x OpenAI

Corporate control & licensed IP

Reactions?

Key Points

- Fully negotiated partnership giving OpenAI limited access to Disney's IP
- Disney kept strict control over how characters and assets could be used
- Aimed to speed up production and support creative teams
- Highlights how powerful companies can adopt AI on their own terms
- Raises questions about who counts as a creator when AI is involved

Case 2: Andersen v. Stability AI
*Independent creators & data
scraping*

Reactions?

Key Points

- Artists allege their work was scraped without consent to train AI models
- Models can mimic their styles, affecting income and recognition
- Companies claim training is “fair use”; courts have not decided
- Independent creators lack bargaining power or control
- Shows how AI systems rely on uncompensated creative labor

What these cases reveal...

Across both cases, we see that:

- Power determines who can negotiate how their work is used
- AI systems rely on human creative labor, but credit and compensation are uneven
- Authorship becomes less about making and more about control
- Consent looks very different for corporations vs. individual creators

Our Take on AI Authorship

AI cannot meaningfully be considered an author in its current form.

No intention

AI has no goal, desire, or reason behind what it makes.

No judgement

AI cannot decide what matters or what is meaningful.

No accountability

AI cannot be held responsible for what it produces.

A framework moving forward

Authorship in the age of AI is not a settled question, it is a constant ongoing negotiation.



Transparency

Clearly disclose how AI systems are built and what data was used to train them.



Attribution

Credit the human roles involved. Dataset creators, prompt writers, curators, and etc.



Protection

Strengthen protections for those whose work and labor make these systems possible.

References...

Ahuja, V. K. (2020, June 11). Artificial Intelligence and Copyright: Issues and Challenges. *SSRN*.
https://papers.ssrn.com/sol3/papers.cfm?abstract_id=3864922

Birhane, A. (2021). Algorithmic injustice: A relational ethics approach. *Patterns (New York, N.Y.)*, 2(2), 100205.
<https://doi.org/10.1016/j.patter.2021.100205>

Boden, M. A. (2014). Creativity and artificial intelligence. In *The philosophy of creativity*. Oxford University Press.
<https://doi.org/10.1093/acprof:oso/9780199836963.003.0012>

Coeckelbergh, M. (2013). David J. Gunkel: *The machine question: Critical perspectives on AI, robots, and ethics*. Springer Netherlands. <https://doi.org/10.1007/s10676-012-9305-y>

Crawford, J., Cowling, M., Ashton-Hay, S., Kelder, J.-A., & Middleton, R. (2023). Artificial intelligence and authorship editor policy: ChatGPT, Bard, Bing AI, and beyond. *Journal of University Teaching and Learning Practice*, 20(5), 1–11. <https://search.informit.org/doi/10.3316/informit.T2024112700020401244487741>

Moulaison-Sandy, H. (2023). What is a person? Emerging interpretations of AI authorship and attribution. *Proceedings of the Association for Information Science and Technology*, 60(1), 279–290.
<https://doi.org/10.1002/pr2.788>

... continued

Pasquinelli, M., & Joler, V. (2021). The nooscope manifested: AI as instrument of knowledge extractivism. *AI & Society*, 36(4), 1263–1280. <https://doi.org/10.1007/s00146-020-01097-6>

Samuelson, P. (2023). Generative AI meets copyright. *Science*. <https://doi.org/10.1126/science.adi0656>

Terranova, T. (2000). Free labor: Producing culture for the digital economy. *Social Text*, 18(2), 33–58. https://doi.org/10.1215/01642472-18-2_63-33

Thaler v. Perlmutter, 687 F. Supp. 3d 140 (D.D.C. 2023), aff'd, No. 23-5233 (D.C. Cir. 2025). <https://fairuse.stanford.edu/case/thaler-v-perlmutter/>

U.S. Copyright Office. (2023, March 16). Copyright registration guidance: Works containing material generated by artificial intelligence (88 Fed. Reg. 16,190). U.S. Copyright Office. https://www.copyright.gov/ai/ai_policy_guidance.pdf

U.S. Copyright Office. (2025, January 29). Copyright and artificial intelligence, Part 2: Copyrightability. U.S. Copyright Office. <https://www.copyright.gov/ai/Copyright-and-Artificial-Intelligence-Part-2-Copyrightability-Report.pdf>

Yang, S. A., & Zhang, A. H. (2024). Generative AI and copyright: A dynamic perspective. *arXiv*. <https://arxiv.org/abs/2402.17801>

Thank You!

Questions?

Dear EDM,

Thank you for a thought-provoking, engaging class session!

In terms of the presentation components of your session, the general structure of your presentation was clear and coherent. It is helpful to have a consistent design template for all the slides to convey a sense of unity in the presentation; this also provides a more polished impression. (It is also helpful to include citations for your references on the slides. Great!) Likewise, it is excellent to provide an outline of what you will talk about. I also appreciated the intentional grounding provided by your guiding question: “What does authorship mean in a world of generative AI?”

Your topic requires the communication of different kinds of complicated ideas, with philosophical, technical, and legal components. Your presentation managed this very well, particularly in the first section, which introduced the notion of creativity and described how generative AI challenges our conceptions. Boden’s definition of creativity provided useful grounding for this discussion.

In presenting the two examples of images for the class to consider, asking if an image “feels original” was an interesting proposition, because this encourages us to think about the experience of creation as well as the output of creation. As one of your classmates noted at one point, “If Max created that image, how awesome you are! But if AI created it, bummer.” This will be a compelling idea to surface in your symposium panel as well.

You condensed an impressive amount of information into your technical overview of AI. It was especially useful here to let the slides carry a lot of the detail, while the oral presentation provided more of a summary. However, at some points, the slides weren’t on the screen long enough for me to fully read them (and I am a pretty fast reader). So you may want to practice with an audience to ensure that you provide enough time for the slides to be fully assimilated by the viewer. (It was easier for me to read all of the slides in the legal overview portion, which also summarized an extensive set of information.)

The way that the presentation emphasized the relationship between authorship and power, including the imbalance of power between the creators of AI models (corporations) and the creators of data to train the models (individuals). This provided an effective framework for the legal association of authorship with ownership, and for the introduction of copyright.

In terms of the interactive components of your session, the initial group brainstorm regarding creativity sparked interest in your subsequent presentation. The case studies were current and substantive, and provided lots to talk about. The lack of time towards the end for a more extended group debrief was unfortunate, but the group discussions were lively.

You ran out of time when it came to presenting the proposals suggested by your group; be sure to practice your symposium panel to ensure that you have sufficient time left for this portion of your material.

Once again, thank you for a great class! Marvelous work!